

MY457/MY557: Causal Inference for Observational and Experimental Studies

Week 11: Difference-in-Differences 2

Daniel de Kadt

Department of Methodology
LSE

Winter Term 2026

Course Outline

- **Week 1:** Causal frameworks
- **Week 2:** Randomization
- **Week 3:** Selection on observables I
- **Week 4:** Selection on observables II
- **Week 5:** Selection on observables III
- *Week 6: Reading week*
- **Week 7:** Instrumental variables I
- **Week 8:** Instrumental variables II
- **Week 9:** Regression discontinuity
- **Week 10:** Difference-in-differences I
- **Week 11:** Difference-in-differences II

Difference-in-Differences So Far

Last week we studied the **canonical 2-period difference-in-differences** (DiD) design, with a brief foray into a special 3-period case.

Identification and estimation was reasonably straightforward:

- Key identification assumption is **parallel trends**, plus **no anticipation**
- Use either a plug-in or regression-based estimator

However, we often encounter DiD settings that are **more complex**:

- More than 2 time periods
- Treatment is assigned variably over time
- Treatment effects are heterogeneous (over e.g. time)
- Treatment is non-binary

Today we will consider **identification and estimation** in such settings.

- 1 Motivating Example
- 2 Fixed Effects
- 3 Variable Treatment Timing
- 4 Multi-Period Designs with Heterogeneous Treatment Effects
- 5 Synthetic Control Method Primer

- 1 Motivating Example
- 2 Fixed Effects
- 3 Variable Treatment Timing
- 4 Multi-Period Designs with Heterogeneous Treatment Effects
- 5 Synthetic Control Method Primer

Minimum Wages and Employment

Do higher minimum wages decrease employment?


Card and Krueger (1994) consider impact of New Jersey's 1992 **minimum wage increase** from \$4.25 to \$5.05 per hour.

Compare employment in 410 fast-food restaurants in New Jersey and eastern Pennsylvania before and after the rise.

Survey data on wages and employment from **two waves**:

- Wave 1: March 1992, one month **before** the minimum wage increase
- Wave 2: December 1992, eight months **after** increase

Minimum Wages and Employment



Minimum Wages and Employment

Variable	Stores by state		
			Difference,
	PA (i)	NJ (ii)	NJ – PA (iii)
1. FTE employment before, all available observations	23.33 (1.35)	20.44 (0.51)	– 2.89 (1.44)
2. FTE employment after, all available observations	21.17 (0.94)	21.03 (0.52)	– 0.14 (1.07)
3. Change in mean FTE employment	– 2.16 (1.25)	0.59 (0.54)	2.76 (1.36)

Treatment: Increase of minimum wage in New Jersey

Difference 1: Pre-minimum wage - Post-minimum wage

Difference 2: New Jersey- Pennsylvania

Minimum Wage and Employment

Card & Krueger (1994) found a **positive effect of minimum wage on employment**.


Card & Krueger (2000) revisited this design and setting with better data, and found **no effect** either way.

Lots of debate followed – many papers reconsidered this question using a more general approach: Leveraging cross- and within-state variation throughout the USA. They largely found **negative effects** on employment.

Dube, Lester, and Reich (2010) revisit this debate:

- Find all cross-state-border changes in MW policies (1990 - 2006)
- Collect earnings and employment data for every county in the USA in this time period.
- Generalize the DiD case study approach.


Variation in Space




Estimated Dynamic Effects – Entire Sample

Ln Earnings

1. All County Sample, Common Period Effects





Ln Employment



Estimated Dynamic Effects – Border Sample Only

5. Contiguous Border County-Pair Sample, Common Period Effects



- 1 Motivating Example
- 2 Fixed Effects
- 3 Variable Treatment Timing
- 4 Multi-Period Designs with Heterogeneous Treatment Effects
- 5 Synthetic Control Method Primer

Fixed Effects Estimation and Difference in Differences

Recall the **additive linear model for panel data** with 2 periods:

$$Y_{it}(z) = \alpha_i + \gamma t + \tau z + \varepsilon_{it}$$

where

- $i \in \{0, \dots, N\}$: Unit indicator
- $t \in \{0, 1\}$: Time indicator
- $Y_{it}(z)$: potential outcome under treatment status $Z \in \{0, 1\}$
- α_i : **time-invariant unit-specific unobserved effect**
- ε_{it} : idiosyncratic error term
- τ : (homogeneous, constant) treatment effect of interest

In a 2-period design, we saw that the first-difference regression:

- Unbiasedly estimates τ under parallel trends and no anticipation
- τ will coincide with τ_{ATE} and τ_{ATT} if assumed model is correct

Fixed effects models generalise this to $t > 2$. We will first consider the **assumptions** required by different models (read: assumed data generating processes), then turn to **estimation**.

Panel Data Notation and Setup

Let's introduce some generalised notation for panel data:

- **Unit**: Distinct analytical entities (e.g. people, countries, schools).
- **Observation**: One 'peek' at one unit (e.g. year, district, student).
- y_{it} : Observed outcome for unit i in period t
- $Z_{it} \in \{0, 1\}$: Treatment indicator for unit i in period t

Note 1: We are using **time** as our dimension of observation, but it could be other things – this goes beyond this lecture!

Note 2: Time-varying **covariates** can be included alongside Z_{it} – we omit them here for simplicity.

In a given panel, we observe a sample of $i = 1, 2, \dots, N$ distinct units at $t = 1, 2, \dots, T$ time periods (a 'balanced panel')

Panel Data Notation and Setup

Collect variables for unit i :

$$\mathbf{y}_i = \begin{pmatrix} y_{i1} \\ \vdots \\ y_{it} \\ \vdots \\ y_{iT} \end{pmatrix}_{T \times 1} \quad \mathbf{z}_i = \begin{pmatrix} z_{i1} \\ \vdots \\ z_{it} \\ \vdots \\ z_{iT} \end{pmatrix}_{T \times 1}$$

And stack them for all units (a 'long panel'):

$$\mathbf{y} = \begin{pmatrix} \mathbf{y}_1 \\ \vdots \\ \mathbf{y}_i \\ \vdots \\ \mathbf{y}_N \end{pmatrix}_{NT \times 1} \quad \mathbf{z} = \begin{pmatrix} \mathbf{z}_1 \\ \vdots \\ \mathbf{z}_i \\ \vdots \\ \mathbf{z}_N \end{pmatrix}_{NT \times 1}$$

Pooled OLS Model

Let's start by ignoring the panel structure entirely, and assume:

$$y_{it} = \tau Z_{it} + v_{it}, \quad t = 1, 2, \dots, T$$

with **composite error** $v_{it} \equiv \alpha_i + \varepsilon_{it}$


We could estimate this model directly using OLS. Our estimate of τ will be unbiased and consistent under **contemporaneous exogeneity**:

$$\mathbb{E}[v_{it}|Z_{it}] = 0 \text{ for } t = 1, 2, \dots, T$$

Read: Since we control for nothing, we identify using **all variation** in Z_{it} . We must assume this variation is ignorable with respect to potential outcomes.

When might this be credible? Treatment is a **shock** that is **as-if random** with respect to **both units and time**.

Pooled OLS: Illustration



Unit Fixed Effects Model

Now consider a **unit fixed effects** assumed model:

$$y_{it} = \tau Z_{it} + \alpha_i + \varepsilon_{it}, \quad t = 1, 2, \dots, T$$

where α_i captures **time-invariant unit differences**

We could estimate this model using OLS. Our estimate of τ will be unbiased and consistent under **strict exogeneity**:

$$\mathbb{E}[\varepsilon_{it} | Z_{i1}, Z_{i2}, \dots, Z_{iT}, \alpha_i] = 0, \quad t = 1, 2, \dots, T$$

Read: Once we control for α_i , we identify using **within-unit over-time variation** in Z_{it} . We assume this variation is (conditionally) ignorable.

When might this be credible? When treatment is a **shock** that is **as-if random** with respect to **time**.

If the assumed model is correct, this is a generalization of the **pre-post design**!

(Unit) Fixed Effects Estimation

Three equivalent approaches to estimating τ in the unit FE model:

1. The **LSDV estimator**:

- a. Regress y_{it} on Z_{it} and unit dummies

Note 1: Here we directly estimate α_j

Note 2: With large N , this can be computationally expensive

2. **Within** estimation:

- a. Create demeaned variables: $\ddot{Z}_{it} \equiv Z_{it} - \bar{Z}_i$ and $\ddot{y}_{it} \equiv y_{it} - \bar{y}_i$

- b. Regress $\ddot{y}_{it}$ on $\ddot{Z}_{it}$

Note: By within-demeaning we 'purge' the fixed effects α_j

3. **First differences** estimation (like last week):


- a. Create differenced variables: $\Delta Z_{it} = Z_{it} - Z_{i,t-1}$ and $\Delta y_{it} = y_{it} - y_{i,t-1}$

- b. Regress Δy_{it} on ΔZ_{it}


Note: First differencing also purges α_j

All three yield the same $\widehat{\tau_{FE}}$.


Within Estimation: Unit Means



Within Estimation: De-meaning



Within Estimation: De-meaned OLS



Adding Time Effects

Consider again our assumed 'unit fixed effects' model:

$$y_{it} = \tau Z_{it} + \alpha_i + \varepsilon_{it}, \quad t = 1, 2, \dots, T$$

Recall we assume on ignorable within-unit temporal variation in Z_{it} .

Typical violation: **common changes over time** that affect all units' y_{it} in the same way and are correlated with Z_{it} (e.g. **history** or **maturation** effects).

More realistic models might include **time effects**:

- linear time trends
- non-linear time trends
- unit-specific time trends
- **time fixed effects**

Two-Way Fixed Effects Regression

Let's add **time fixed effects** to our assumed model:

$$y_{it} = \tau Z_{it} + \alpha_i + \delta_t + \varepsilon_{it}, \quad t = 1, 2, \dots, T$$

where

- α_i captures time-invariant unit differences
- δ_t captures common time-specific differences

This is the **two-way fixed effects** (TWFE) model.

Once we account for α_i and δ_t , we identify using **idiosyncratic variation** in Z_{it} . We assume this variation is ignorable.

When might this be credible? Treatment is a **shock** whose **timing** varies across units for **idiosyncratic** reasons. If the assumed model is correct, our estimate of τ is **generalized difference-in-differences**.

Estimation as before, extending each approach by including time dummies. In R, use `plm` or `fixest` (preferred).

Dynamic Two-Way Fixed Effects

We can specify a TWFE model allowing **dynamic (time-varying)** treatment effects:

$$y_{it} = \alpha_i + \delta_t + \sum_{r \neq 0} \mathbf{1}[R_{it} = r] \tau_r + \varepsilon_{it}$$

where

- α_i captures time-invariant unit differences
- δ_t captures common changes over time
- R_{it} is the period relative to treatment for unit i , with $R_{it} = 0$ denoting the last pre-treatment period (reference category)
- τ_r is a relative-period treatment effect

Read: This **event study** estimator allows for heterogeneous treatment effects of a specific form: **constant τ_r across treatment cohorts**.

- 1 Motivating Example
- 2 Fixed Effects
- 3 Variable Treatment Timing**
- 4 Multi-Period Designs with Heterogeneous Treatment Effects
- 5 Synthetic Control Method Primer

Variable Treatment Timing

Multi-period treatment regimes usually vary over two dimensions:

- **Uniform vs. Staggered**: Does treatment occur simultaneously, or over time?
- **Absorbing vs. Non-absorbing**: Once treatment occurs, can it switch off?

With anything other than uniform and absorbing treatment timing, TWFE for DiD may not behave well. For synthesis, see:

- Baker, Larcker, and Wang (2022)
- Roth, Sant'Anna, Bilinski, and Poe (2023)


Short of further assumptions, the estimand targeted by TWFE is **not easily interpretable** $\rightsquigarrow$ it is a weighted average of many different treatment effects.

These weights can be negative (!), are generally non-intuitive, and can potentially severely mislead (e.g. sign-flips).


$\widehat{\tau}_{TWFE}$ Decomposition

To see this, we can **decompose** $\widehat{\tau}_{TWFE}$. We focus on Goodman-Bacon (2021).

Define three groups: never treated (U), early treated (k), and late treated (l)



τ_{TWFE} Decomposition



τ_{TWFE} is the weighted average of these four 2x2 treatment effects.

$\widehat{\tau}_{TWFE}$ Decomposition

$$\widehat{\tau}_{TWFE} = \sum_{k \neq U} s_{kU} \hat{\tau}_{kU}^{2 \times 2} + \sum_{k \neq U} \sum_{\ell > k} [s_{kl}^k \hat{\tau}_{k\ell}^{2 \times 2, k} + s_{kl}^l \hat{\tau}_{k\ell}^{2 \times 2, \ell}]$$

where

- $\hat{\tau}^{2 \times 2}$ are different 2x2 estimators
- s are estimator-specific weights

The weights are a function of three things:


1. 2x2 subsample size (higher, $\uparrow$ weight)
2. Ratio of treated to control (closer to 0.5, $\uparrow$ weight)
3. Relative timing of treatment (more central, $\uparrow$ weight)

Problem: Some of these comparisons may be **forbidden** such that already-treated units used as controls after they are treated.

Intuition: We subtract off changes in untreated outcomes **and treatment effects**.

τ_{TWE} Decomposition

Weighting is heavily dependent on timing:



Key Assumptions for TWFE as Generalised DiD

This decomposition reveals the assumptions under which **traditional TWFE** might be trusted with **multi-period panel data** and a **DiD** design.

A **causal DiD** interpretation with **TWFE** requires either:

1. Parallel trends, no anticipation, and homogenous τ

Or:

2. Parallel trends, no anticipation, and uniform timing (constant τ over time)

Quickly explore how much of a problem this may be using **bacondecomp** in R.

- 1 Motivating Example
- 2 Fixed Effects
- 3 Variable Treatment Timing
- 4 Multi-Period Designs with Heterogeneous Treatment Effects
- 5 Synthetic Control Method Primer

Two Classes of 'Modern' Estimators

Multi-period DiD with **non-uniform (staggered or non-staggered) treatment timing** should be approached with **caution**.

Two general types of modern estimators that can help:

1. Flexible matching and re-weighting estimators:

- Make the 'right' comparisons only, weight appropriately, and recover τ_{ATT} .
- Many estimators exist: Strehzhev (2018), de Chaisemartin and D'Haultfœuille (2020), Sun and Abraham (2021), Imai and Kim (2021), Dube et al. (2023), de Chaisemartin and D'Haultfœuille (2024)
- We will focus on Callaway and Sant'Anna (2021)

2. Counterfactual estimators:

- Estimate only Y_0 , thus avoiding forbidden comparisons, and recover τ_{ATT} .
- Many estimators exist: Gobillon and Magnac (2016), Xu (2017), Borusyak et al. (2021), Gardner (2021), Wooldridge (2021)
- We will focus on Liu, Wang, and Xu (2022)

Callaway and Sant'Anna (2021): Setup

Callaway and Sant'Anna (2021) study a DiD setting with **multiple time periods**, **staggered** treatment timing, there may be **heterogeneous** τ , and parallel trends may hold only conditional on $\mathbf{X}$.

They begin by defining a new estimand, the **group-time ATT**:

$$\tau_{g,t}^{ATT} = \mathbb{E}[Y_t(1) - Y_t(0) | G_g = 1]$$

where

- there are $T = t \in \{1, \dots, T\}$ time periods
- $G_g \in \{0, 1\}$ indicates whether a unit is first treated in period g
- $Y_t(1)$ and $Y_t(0)$ are potential outcomes under treatment and control, at $T = t$

Intuitively, this has already brought us a long way. We can now reason in abstraction about **every** 2x2 comparison in our data.

Callaway and Sant'Anna (2021): Identification

If **parallel trends** holds, this group-time ATT can be identified as:

$$\tau_{g,t}^{ATT} = \mathbb{E}[\underbrace{Y_t - Y_{g-1}}_{\text{Long difference}} | G_g = 1] - \mathbb{E}[\underbrace{Y_t - Y_{g-1}}_{\text{Long difference}} | C = 1]$$

where $C \in \{0, 1\}$, taking 1 if never treated (no forbidden comparisons!)

Read: For any treated cohort, at any time period, conduct a diff-in-diff where:

- Difference 1: $T = t$ vs. $T = g - 1$ (long difference)
- Difference 2: G_g (cohort) vs. C (never treated)

Long difference in Y : change in Y between any time period t and the **last pre-treatment period**, $g - 1$.

Callaway and Sant'Anna (2021): Identification

If **conditional parallel trends** holds:

$$\tau_{g,t}^{ATT} = \mathbb{E} \left[\underbrace{\left(\frac{G_g}{\mathbb{E}[G_g]} - \frac{\frac{p_g(X)C}{1-p_g(X)}}{\mathbb{E} \left[\frac{p_g(X)C}{1-p_g(X)} \right]} \right)}_{\text{Weights}} \underbrace{(Y_t - Y_{g-1})}_{\text{Long difference}} \right]$$

$p_g(X) = P(G_g = 1|X, G_g + C = 1)$ is a propensity score

Read: Up-weight control units where $p_g(X)$ is similar to the **group-specific treated units**

Callaway and Sant'Anna (2021): Estimation


Estimation proceeds as follows:

1. Estimate $\hat{p}_g$ for each group g
2. Estimate $\hat{\tau}_{g,t}^{ATT}$ by plugging in fitted values and observed Y into the (estimator-version) of the expression on the previous slide
3. Combine the estimated values of $\hat{\tau}_{g,t}^{ATT}$ to retrieve quantities of interest

Some quantities of interest:

- Simple average of $\hat{\tau}_{g,t}^{ATT}$ across t and g
- Weighted average of $\hat{\tau}_{g,t}^{ATT}$ weighting by group sizes
- Any other principled summary measure!

Callaway and Sant'Anna (2021) on the Minimum Wage



Liu, Wang, and Xu (2022): Setup

Liu, Wang, and Xu (2022) consider DiD settings with **multiple** time periods, **staggered** treatment timing that **may or may not be absorbing**, and there may be **heterogeneous** τ .

Define the **estimand** of interest as:

$$\tau_{ATT} = \mathbb{E}[Y_{it}(1) - Y_{it}(0) \mid Z_{it} = 1, C_i = 1]$$

where

- Z_{it} is our DiD treatment indicator
- C_i is an indicator for 'ever treated' units
- $Y_{it}(1)$ and $Y_{it}(0)$ are potential outcomes under treatment and control

Idea: Estimate **only** $Y_{it}(0)$ using pre-treatment data, taking $Y_{it}(1)$ as missing.

Estimate τ_{ATT} by taking differences between $Y(1)$ and $Y(\hat{0})$.

Note: This is a philosophical departure from TWFE!

Liu, Wang, and Xu (2022): Estimation

Authors offer three estimators:

- FEct Estimator:

$$Y_{it}(0) = \mathbf{x}_{it}^{\top} \boldsymbol{\beta} + \alpha_i + \delta_t + \varepsilon_{it}$$

- IFect Estimator:

$$Y_{it}(0) = \mathbf{x}_{it}^{\top} \boldsymbol{\beta} + \alpha_i + \delta_t + \lambda_i' \mathbf{f}_t + \varepsilon_{it}$$

- MC Estimator:

$$\mathbf{Y}(0) = \mathbf{x}_{it}^{\top} \boldsymbol{\beta} + \mathbf{L} + \boldsymbol{\varepsilon}$$

Liu, Wang, and Xu (2022): Interactive Fixed Effects

Let's look briefly at the **interactive fixed effects** estimator.

Within the TWFE setup, replace $\alpha_i + \delta_t$ with $\alpha_i + \delta_t + \lambda_i^\top \mathbf{f}_t$ (Bai, 2009):


- $\mathbf{f}_t$: unobserved **common factors** – time-varying, shared by all units
- λ_i : unit-specific **factor loadings** – how sensitive unit i is to each factor

Read: If parallel trends fails because time effects are not actually common, IFECt finds those differences in the pre-treatment data and adjusts the counterfactual.

Note: This does not require covariates (though they can be used). Instead amounts to finding a low-dimensional representation of $\mathbf{Y}(\mathbf{0})$ in the pre-period.

Estimation uses an **EM algorithm** on untreated observations to jointly recover $\hat{\mathbf{f}}_t$ and $\hat{\lambda}_i$; the number of factors r is selected via **cross-validation**.

Liu, Wang, and Xu (2022): Simulated Example



Some Practical Advice

For computation, most of these routines are well packaged in most software:

1. For Callaway & Sant'Anna package `{did}` in R.
2. For Liu et al, package `{fect}` in R.
3. See Asjad Naqvi's summary: <https://asjadnaqvi.github.io/DiD/>

Be aware, however:

- TWFE is often okay! Try a number of estimators, report honestly.
- 'Modern' estimators have low power (see Chiu et al, 2026, and Weiss, 2026)
- Low power may mean quite variable point estimates.
- Focus on **research design and falsification** first!

Design precedes estimation!

- 1 Motivating Example
- 2 Fixed Effects
- 3 Variable Treatment Timing
- 4 Multi-Period Designs with Heterogeneous Treatment Effects
- 5 Synthetic Control Method Primer

Synthetic Control Method: Primer

DiD requires parallel trends in the expected value of potential outcomes.
Generally cannot help where there are **time- and unit-varying confounders**.

Synthetic Control Methods (SCM) take a **different approach**:

1. Find $\mathbf{W}^*$, an $N - 1$ length vector of unit weights that minimizes $||\mathbf{X}_1 - \mathbf{X}_0 \mathbf{w}||$ for $\mathbf{X}_1$ a matrix of pre-treatment **outcomes and covariates** for the treated unit, and $\mathbf{X}_0$ likewise for the control.

Read: Find the weighted combination of control units that best matches – in **both levels and trend** – the treated unit in the pre-treatment period.

This weighted set of control units is the **synthetic control**.


2. An approximately unbiased estimator of the unit-specific effect in the post-treatment period is:

$$\hat{\tau}_{1t} = Y_{1t} - \sum_{j=2}^{J+1} w_j^* Y_{jt} \quad \text{for } t \in \{T_0 + 1, \dots, T\}$$


3. Inference uses placebos (roughly, randomization inference).

For primer on use, see Abadie (2020).


SCM Example: California Prop 99 (Abadie et al. 2010)



SCM Example: California Prop 99 (Abadie et al. 2010)



SCM Example: California Prop 99 (Abadie et al. 2010)





"Be very very quiet, we're hunting identifying variation in D "

Temporary page!

$\text{\LaTeX}$ was unable to guess the total number of pages correctly. As there was unprocessed data that should have been added to the final page this extra page has been added to receive it.

If you rerun the document (without altering it) this surplus page will go away because $\text{\LaTeX}$ now knows how many pages to expect for this document.